

LKP-tests

Linux kernel performance test tool

What is LKP-tests

- Framework to run benchmarks
 - Fetch/install the benchmarks
 - Prepare test environment
 - For example, format disk
 - Run benchmark with various parameters
 - Collect performance statistics
 - Benchmark/performance statistics parsing
- Framework for performance analysis
 - Result compare
- Used for benchmark running and reproducing regression captured by 0-Day

Highlights

- Rich benchmarks support
- Capture performance statistics (monitors) in addition to benchmark score
 - For example, perf stat, perf profile, vmstat, etc.
- Compare/analyze all performance statistics data (benchmark and monitors)
 - Unified data model

Benchmarks

- Integrated ~50 benchmarks
 - Fetch/install benchmarks
 - Run benchmarks with various parameters
 - Parse result of benchmarks
- Mostly micro-benchmarks, several macro-benchmarks
- Mostly server workload

Benchmark List - 1

- Scheduler
 - Hackbench, perf-bench-sched-pipe, etc.
- File system/IO
 - Fio, fsmark, iozone, dbench, dd, fileio, postmark, etc.
- Network
 - Netperf, iperf, apachebench, ku-latency, nepim, netpipe, nuttcp, qperf, siege, sockperf, stutter, tbench, thrulay, etc.
- Scalability
 - Reaim, aim7, will-it-scale, etc.
- Memory management
 - vm-scalability, chromeswap, exit_free, pft, perf-bench-numa-mem, pmbench, swapin, tlbflush, etc.

Benchmark List - 2


- Database
 - Pgbench, oltp, etc.
- Noise
 - Ftq, fwq, etc.
- HPC
 - Hpcc, linpack
- Workload emulation
 - Blogbench, ebizzy (web application server), kbuild, pbzip2, pigz, pixz, plzip, pxz, etc.
- Others
 - Tcrypt, unixbench

Functionality Test

- Memory Management
 - Bust_shm_exit, libhugetlbfs, ndctl, nvml
- File system
 - Xfstests, ext4-frags, ocfs2test
- Virtualization
 - Kvm-unit-tests
- Network
 - packetdrill
- Others
 - Cpu-hotplug, ftrace_onoff, hwsim(wifi), kernel_selftests, locktorture, ltp, mce-log, mce-test, perf_event_tests, perf-sanity-tests, pm-qa, rcutorture, stress-ng, suspend, test_bpf, trinity, piglit, etc.

Test Coverage

0-Day CI: Test Coverage



Monitors

- Capture performance statistics when benchmark running
 - Important for performance analysis, regression root causing
- About 38 monitors now
- Examples:
 - Perf profile: find the hot spot
 - Vmstat: CPU usage, IO bandwidth, memory usage
 - iostat: IO bandwidth, requests
 - Latencytop: measure latency
 - Turbostat: idle, frequency, and power of the CPU
 - /proc files: meminfo, sched-debug, interrupts, etc.

Setups

- Scripts to prepare the test environment
- About 20 setup scripts
- Example setup scripts
 - disk: wait/check test disk to be ready
 - fs: format disk with file system, mount it
 - cpufreq_governor: CPU frequency governor
 - iosched: choose IO scheduler
 - dirty_thresh: writeback dirty threshold

Jobs

- Define the performance tests
- Benchmark + monitors + setups
 - Each with parameters
- In yaml format
 - Key -> value with array
- Combination of various test parameters
 - One jobfile can define a set of tests
 - For example: ext4+performance, ext4+powersave, xfs+performance, xfs+powersave

fs:

- ext4
- xfs

cpufreq_governor

- performance
- powersave

Job file example

```
# swap-test.yaml
testcase: vm-scalability
swap:
vmstat:
perf-profile:
  delay: 20
# ...
vm-scalability
  test: swap-w-seq
  nr_task:
    - 1
    - 2
    - 4
    - 8
```

Supported distribution

- Debian sid
- Ubuntu 14.04
- Centos 7
- Archlinux

Getting started

```
$ git clone https://github.com/01org/lkp-tests  
$ cd lkp-tests  
$ make install  
$ lkp help
```

Prepare host description

- `lkp-tests/hosts/<hostname>`
 - `Yaml format`
 - `Setup disks for test`
- `Example`

`memory: 128G`

`ssd_partitions: /dev/sda1`

Install packages for a job

```
# browse and select a job you want to run, for example, job/hackbench.yaml  
$ ls lkp-tests/jobs  
$ lkp install $LKP_SRC/jobs/hackbench.yaml
```


Split job

```
# Generate atomic job from one job file
$ lkp split-job lkp-tests/jobs/hackbench.yaml
jobs/hackbench.yaml => ./hackbench-1600%-process-pipe.yaml
jobs/hackbench.yaml => ./hackbench-1600%-process-socket.yaml
jobs/hackbench.yaml => ./hackbench-1600%-threads-pipe.yaml
jobs/hackbench.yaml => ./hackbench-1600%-threads-socket.yaml
jobs/hackbench.yaml => ./hackbench-50%-process-pipe.yaml
jobs/hackbench.yaml => ./hackbench-50%-process-socket.yaml
jobs/hackbench.yaml => ./hackbench-50%-threads-pipe.yaml
jobs/hackbench.yaml => ./hackbench-50%-threads-socket.yaml
```

Run atomic job

```
$ lkp run hackbench-50%-process-pipe.yaml
```

```
...
```

```
2016-11-18 07:55:12 /usr/bin/hackbench -g 1 --process --pipe -l 60000
```

```
Running in process mode with 1 groups using 40 file descriptors each (== 40  
tasks)
```

```
Each sender will pass 60000 messages of 100 bytes
```

```
Time: 7.388
```

```
...
```

Check result and analyze

```
# link to result directory of latest run is created automatically
$ ls result
...
$ lkp result hackbench
/result/hackbench/50%-process-pipe-performance/lkp-kvm/debian/x86_64-rhel-7.
2/gcc-6/4.5.0-2-amd64/1/
/result/hackbench/50%-process-pipe-performance/lkp-kvm/debian/x86_64-rhel-7.
2/gcc-6/4.8.0-1-amd64/0/
$ lkp ncompare -s commit=<kernel1> -o commit=<kernel2>
```

Linux kernel source

- Optional
- Put it in `/c/repo/linux` or `$GIT_ROOT_DIR/linux`

Run user supplied workload

- Use “mytest” test

Gaps

- Bug fixes
- Feedback from more users

References

- Git repo: <https://github.com/01org/lkp-tests>
- Documents in source code: `find lkp-tests -name README.md`
- Intel internal mailing list: lkp@eclists.intel.com
- External mailing list: lkp@01.org

Thanks!

Backup

Cluster: Multiple machine test

- One server and multiple clients
- Steps
 - Queue job for server
 - When schedule server job, queue job for client
 - Sync among server and clients
 - Server and client will all run the cluster job
 - Server run server side
 - Clients run client side

Cluster: Example job file

```
suite: netperf  
testcase: netperf  
category: benchmark
```

```
ip: ipv4
```

```
runtime: 300s  
nr_threads: 200%
```

```
cluster: cs-localhost
```

```
if role server:  
    netserver:
```

```
if role client:  
    netperf:  
        test: TCP_RR
```